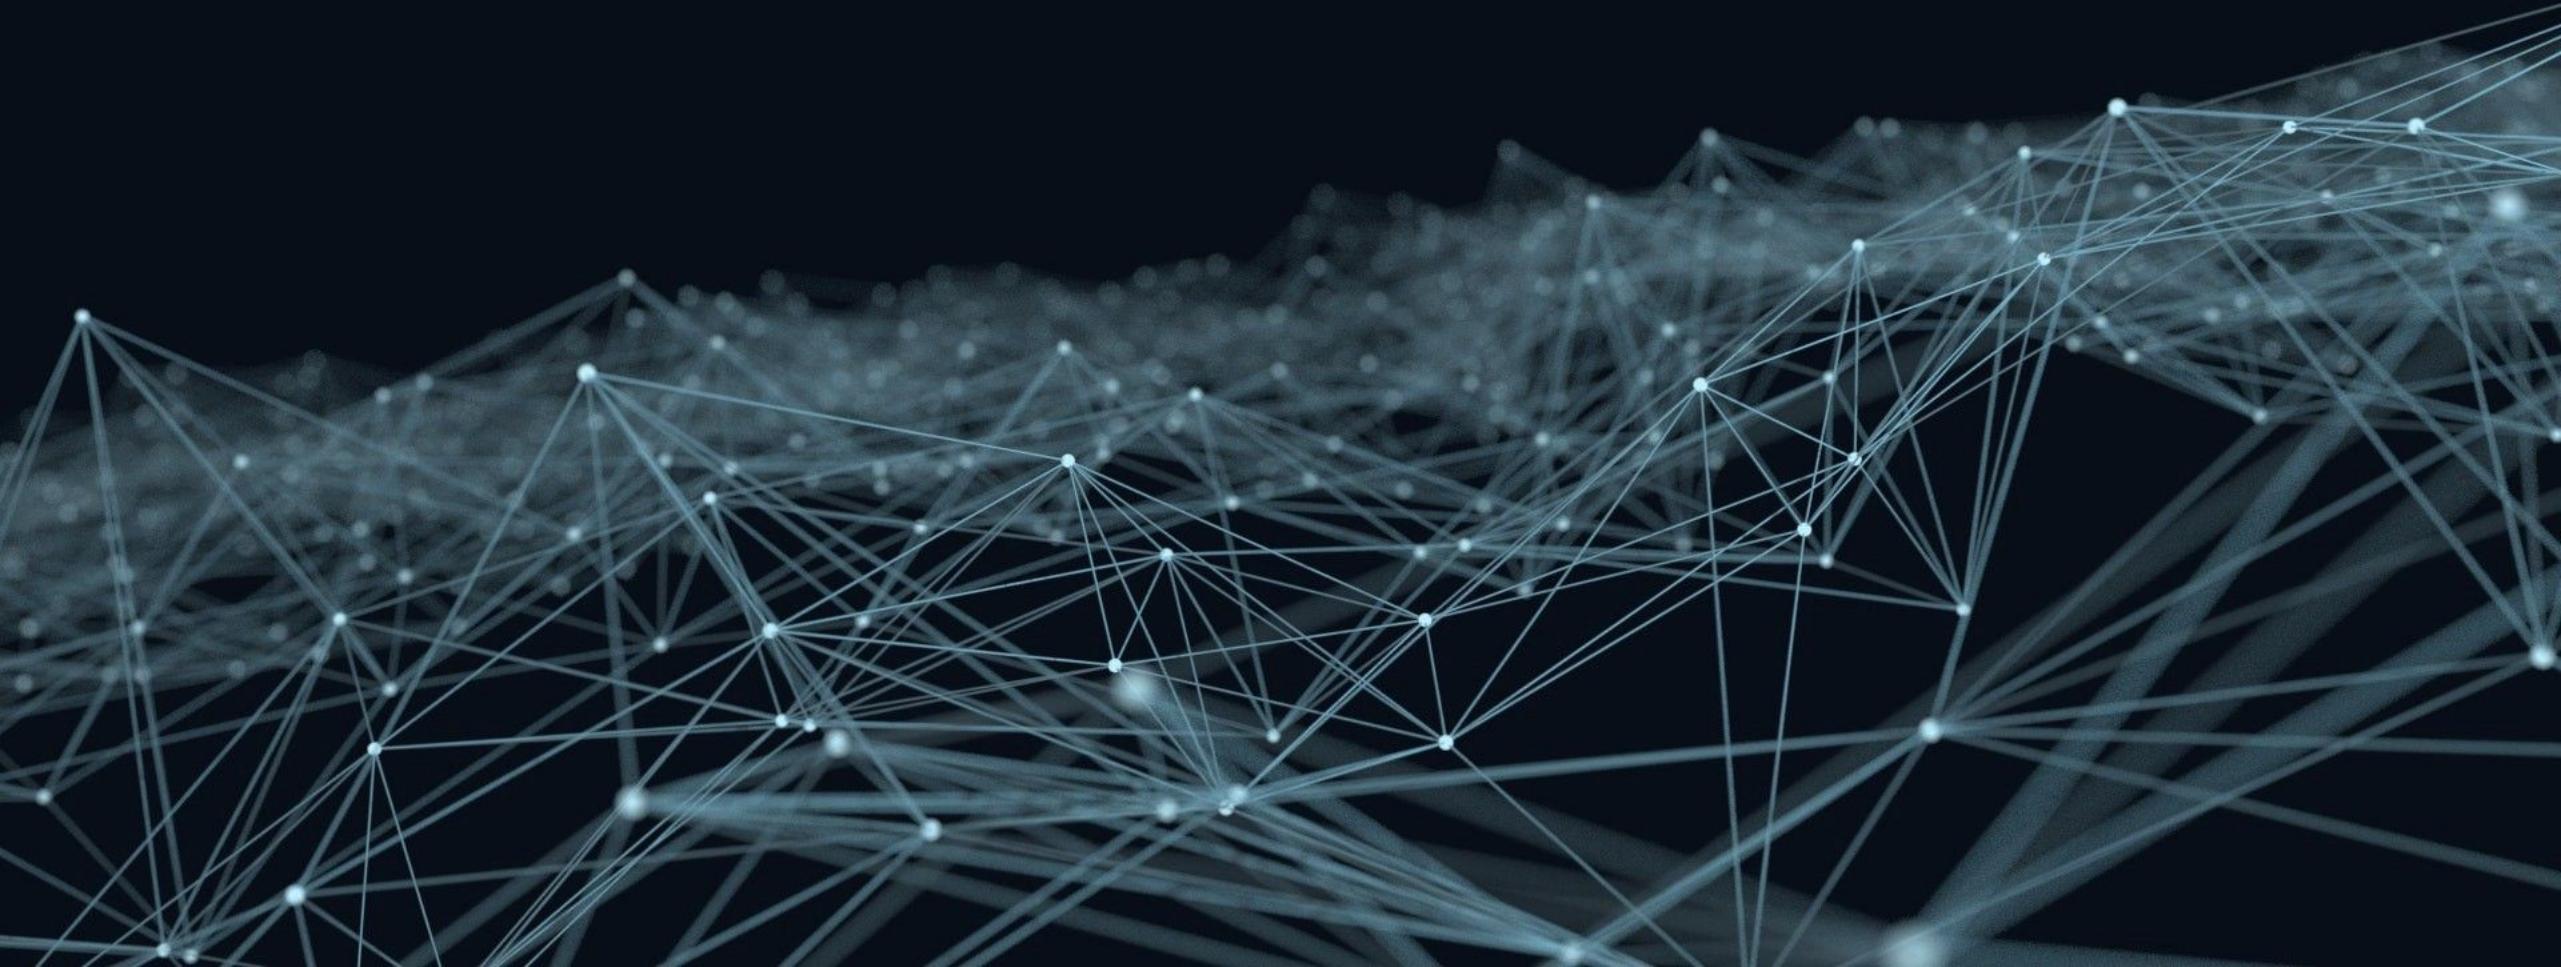


Introduction to CI/CD

Module 5: Jenkins



Origins

- In the early 2000s, most software teams
 - Integrated code infrequently
 - Ran builds manually
 - Discovered integration problems late in the cycle
- In 2004 by Kohsuke Kawaguchi at Sun Microsystems created the Hudson project
 - Open-source written in Java
 - Designed to automate builds and tests
 - Integrated well with Java build tools like Ant and Maven
 - Easy to set up and extensible
 - Focused on automation rather than process enforcement



Hudson to Jenkins (2011)

- In 2010–2011, Oracle acquired Sun Microsystems
 - This led to disagreements between the Hudson community and Oracle regarding:
 - *Project governance*
 - *Trademark ownership*
 - *Contribution control*
- As a result
 - The community forked Hudson
 - The new project was named Jenkins
 - Most contributors and users moved to Jenkins
- From this point on
 - Jenkins became the community-driven project
 - Hudson gradually declined in relevance



Job-Centric CI (2011-2014)

- Early versions of Jenkins focused on:
 - Freestyle jobs
 - UI-driven configuration
 - Scripted build steps
- Characteristics of early Jenkins
 - Configuration stored on the server
 - Pipelines defined through the web UI
 - Heavy reliance on plugins
 - Less reproducible builds
- This reflected the state of CI at the time
 - Automation existed, but IaC was not common



Job-Centric CI (2011-2014)

- Early versions of Jenkins focused on:
 - Freestyle jobs
 - UI-driven configuration
 - Scripted build steps
- Characteristics of early Jenkins
 - Configuration stored on the server
 - Pipelines defined through the web UI
 - Heavy reliance on plugins
 - Less reproducible builds
- This reflected the state of CI at the time
 - Automation existed, but IaC was not common



Pipeline as Code (2014-2016)

- As DevOps practices matured, teams demanded
 - Version-controlled build definitions
 - Reproducibility
 - Better support for complex workflows
- This led to the introduction of the Jenkins Pipeline and Jenkinsfile
- Innovations
 - Pipelines defined in code
 - Stored in Git repositories
 - Reviewed like application code
 - Support for stages, parallelism, and conditions
- This marked a major evolution in Jenkins' design philosophy



DevOps Era (2016-Present)

- Jenkins evolved to support
 - Microservices
 - Containers and Docker
 - Kubernetes-based build agents
 - Cloud-native workflows
- Notable Trends
 - Declarative pipelines simplified syntax
 - Dynamic agents replaced static build servers
 - Integration with modern SCM platforms (GitHub, GitLab)
- Jenkins shifted from being “a CI tool” to a general-purpose automation engine



Jenkins Today

- Why Jenkins is still relevant
 - Extreme flexibility
 - Large plugin ecosystem
 - Strong support for legacy and complex systems
- Challenges
 - Requires operational effort
 - Plugin compatibility management
 - Competition from managed CI/CD platforms
- Despite newer tools, Jenkins remains widely used in:
 - Enterprises
 - Regulated environments
 - Complex, heterogeneous systems



GitLab and GitHub

- GitHub Actions and GitLab CI/CD are modern alternatives to Jenkins' pipelines
 - These are integrated pipeline platforms built directly into their Git hosting services
- These repository tools treat CI/CD as a native feature of the repository
 - Jenkins = external automation engine
 - GitHub/GitLab = pipelines embedded in the SCM platform
- The same basic pipeline concepts apply to these the same as Jenkins
 - The syntax and format of the configuration files
- These turn GitHub and GitLab into orchestration engines



GitHub Actions

- GitHub Actions is GitHub's built-in automation and CI/CD platform
 - Pipelines are defined using YAML workflow files stored in the repository.
- Key characteristics
 - Tight integration with GitHub repositories
 - Event-driven (push, pull request, tag, issue events)
 - Pipelines live in .github/workflows
 - Uses hosted runners or self-hosted runners
- Conceptual flow
 - GitHub event → Action workflow → Jobs → Steps
- Strengths
 - Zero infrastructure to manage
 - Simple setup for GitHub users
 - Strong marketplace of reusable actions



GitLab CI/CD

- GitLab CI/CD is a first-class feature of GitLab, not an add-on.
 - Pipelines are defined in a single file: “.gitlab-ci.yml”
- Key characteristics
 - CI/CD tightly integrated with GitLab repos
 - Built-in artifact storage and environment tracking
 - Native support for Docker and Kubernetes
 - GitLab Runners execute jobs
- Conceptual flow
 - Git commit → GitLab pipeline → Stages → Jobs
- Strengths
 - Very cohesive developer experience
 - Strong DevOps lifecycle integration
 - Minimal configuration overhead



Comparison

- Configuration style contrast
- Jenkins
 - Highly flexible
 - Scriptable pipelines (Groovy)
 - Powerful but verbose
 - Requires plugin management
- GitHub / GitLab
 - Opinionated YAML syntax
 - Simpler mental model
 - Less customization, faster onboarding
 - Platform constraints apply



Comparison

- Jenkins is often chosen when
 - You need extreme customization
 - You support many SCM systems
 - You have complex legacy workflows
 - You control infrastructure
- GitHub / GitLab pipelines are often chosen when
 - Repositories are already hosted on the platform
 - Teams want fast setup
 - Minimal operational overhead is desired
 - Standard CI/CD patterns are sufficient



Jenkins Distributed Architecture

- The main node is the default node
 - Runs the Jenkins instance
 - Manages access to Jenkins Interface
 - Should be the only node accessible via web interface
 - The master/main node is the controller
 - Best practice is that CICD jobs are delegated to other nodes
- Agent/Slave nodes
 - Tasked with running parts of the CICD jobs
 - Agents are tagged so they can be used selectively
 - Jenkins does not have to be installed on an agent
 - Just Java to run the Jenkins agent jar file
 - Communication takes place via ssh



Jenkins Pipelines

- A pipeline defines a CICD process for that project
- Stages
 - Sequential series of steps to be executed, for example
 - Build > Test > Package > Deploy
- JenkinsFile (infrastructure as code)
 - A script that defines the stages of a pipeline
 - Older form is written in the Groovy scripting language
 - New form is a declarative form of scripting
- Three basic configurations
 - The JenkinsFile is kept in Jenkins
 - The JenkinsFile is kept in a SCM repository
 - The JenkinsFile is kept in the project itself
 - *This allows for multi branch builds*



Post Build Steps

- In addition to stages, Jenkins has a post build stage
- Contains any of a number of clause types
 - Always – always executes
 - Failure – executes only when the pipeline fails
 - Success – executes only when the build succeeds
 - Cleanup – always runs after all the other clauses run
 - Changed – runs if the pipeline or a stage completion status is different from a previous run
- There are more clause types at:
 - <https://jenkins.io/doc/book/pipeline/syntax/#post>



Environment Variables

- There are a number of predefined environment variables
 - Accessed via the global variable “env”
 - Value is accessed using Groovy syntax
 - *“Build ID is \${BUILD_ID}”*
 - We can define environment variables in either the whole pipeline or a given stage using the environment block




Pipeline Features

- Can support complex, real-world, CD Pipeline requirements
 - Pipelines can fork/join, loop, parallel, and more
- Resilient:
 - Pipeline executions can survive master restarts
- Can be paused:
 - Pipelines can pause and wait for human input/approval.
- Pipeline configuration is treated as code in source control



Pipeline



Pipeline Stage View

billing-rest - Stage View

| | Declarative: Checkout SCM | Initialize | Checkout | Build | Publish Reports | SonarQube analysis | ArchiveArtifact | Docker Tag & Push | Deploy - CI | Deploy - QA | Deploy - UAT | Deploy - Production | Declarative: Post Actions |
|--|---------------------------|------------|----------|-------|-----------------|--------------------|-----------------|-------------------|-------------|-------------|--------------|---------------------|---------------------------|
| | 768ms | 1s | 799ms | 57s | 5s | 14s | 125ms | 23s | 37ms | 32ms | 32ms | 32ms | 792ms |
| Average stage times: (Average full run time: ~2min 59s) | | | | | | | | | | | | | |
| #118 | Jul 17 20:58 | 1 commit | ① | 817ms | 1s | 690ms | 1min 36s | 10s | 24s | 198ms | 38s | 38ms | 1s |
| #117 | Jul 16 15:28 | 1 commit | ① | 792ms | 1s | 708ms | 1min 36s | 8s | 23s | 179ms | 38s | 42ms | 2s |
| #116 | Jul 15 21:13 | No Changes | ① | 672ms | 869ms | 694ms | 1min 33s | 10s | 24s | 183ms | 37s | 40ms | 75ms |

Pipeline Features

- Can support complex, real-world, CD Pipeline requirements
 - Pipelines can fork/join, loop, parallel, and more
- Resilient:
 - Pipeline executions can survive master restarts
- Can be paused:
 - Pipelines can pause and wait for human input/approval.
- Pipeline configuration is treated as code in source control



Questions

